

```
import pandas as pd
```

```
df=pd.read_csv('/content/Housing.csv')
print(df.head())
print(df.isnull().sum())
```

```

↗
   price  area  bedrooms  bathrooms  stories  mainroad  guestroom  basement  \
0  13300000  7420         4          2        3        yes         no         no
1  12250000  8960         4          4        4        yes         no         no
2  12250000  9960         3          2        2        yes         no         yes
3  12215000  7500         4          2        2        yes         no         yes
4  11410000  7420         4          1        2        yes         yes        yes

   hotwaterheating  airconditioning  parking  prefarea  furnishingstatus
0                no                yes        2        yes        furnished
1                no                yes        3         no        furnished
2                no                no        2        yes    semi-furnished
3                no                yes        3        yes        furnished
4                no                yes        2         no        furnished
price              0
area              0
bedrooms          0
bathrooms         0
stories           0
mainroad          0
guestroom         0
basement          0
hotwaterheating   0
airconditioning   0
parking           0
prefarea          0
furnishingstatus  0
dtype: int64

```

```

import pandas as pd
df=pd.read_csv('/content/Housing.csv')
missing = df.isnull().sum()
print("Missing values per column:\n", missing[missing > 0])
df = df.dropna()

duplicates = df.duplicated().sum()
print(f'Duplicate rows: {duplicates}')
df = df.drop_duplicates()
df = df[df['price'] < df['price'].quantile(0.99)]
df = df[df['area'] < df['area'].quantile(0.99)]
X = df.drop('price', axis=1)
y = df['price']
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```

↗ Missing values per column:
Series([], dtype: int64)
Duplicate rows: 0

```

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

plt.rcParams['figure.figsize'] = (10, 5)
df = pd.read_csv('/content/Housing.csv')
print("Dataset Shape:", df.shape)
print("\nData Types and Nulls:\n", df.info())

```

```

Dataset Shape: (545, 13)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 545 non-null   int64
1   area                 545 non-null   int64
2   bedrooms             545 non-null   int64
3   bathrooms            545 non-null   int64
4   stories              545 non-null   int64
5   mainroad             545 non-null   object
6   guestroom            545 non-null   object
7   basement             545 non-null   object
8   hotwaterheating      545 non-null   object
9   airconditioning      545 non-null   object
10  parking              545 non-null   int64
11  prefarea             545 non-null   object
12  furnishingstatus     545 non-null   object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB

```

Data Types and Nulls:
None

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
print("Dataset Shape:", df.shape)
print(df.info())
print(df.describe())

```

```

Dataset Shape: (545, 13)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 545 non-null   int64
1   area                 545 non-null   int64
2   bedrooms             545 non-null   int64
3   bathrooms            545 non-null   int64
4   stories              545 non-null   int64
5   mainroad             545 non-null   object
6   guestroom            545 non-null   object
7   basement             545 non-null   object
8   hotwaterheating      545 non-null   object
9   airconditioning      545 non-null   object
10  parking              545 non-null   int64
11  prefarea             545 non-null   object
12  furnishingstatus     545 non-null   object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
None

```

	price	area	bedrooms	bathrooms	stories \
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000

	parking
count	545.000000
mean	0.693578
std	0.861586
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	3.000000

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv('/content/Housing.csv')
sns.set(style="whitegrid")
plt.rcParams['figure.figsize'] = (10, 4)
sns.histplot(df['price'], kde=True, bins=30)
plt.title('Sale Price Distribution')
plt.xlabel('SalePrice')
plt.show()



```




```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv('/content/Housing.csv')
sns.set(style="whitegrid")
plt.rcParams['figure.figsize'] = (9, 4)
categorical_features = ['bedrooms', 'bathrooms']
for feature in categorical_features:
    sns.boxplot(x=df[feature], y=df['price'])
    plt.title(f'{feature} vs price')
    plt.xticks(rotation=45)
plt.show()

```



```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv('/content/Housing.csv')
sns.set(style="whitegrid")
plt.rcParams['figure.figsize'] = (9, 4)
numerical_features = ['area', 'basement']
for feature in numerical_features:
    sns.scatterplot(x=df[feature], y=df['price'])
    plt.title(f'{feature} vs price')
    plt.xlabel(feature)
    plt.ylabel('SalePrice')
    plt.show()
```



```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

df = pd.read_csv("/content/Housing.csv")

print(df.head())
print(df.info())

df.columns = df.columns.str.lower()

df = df.dropna()

categorical_features = df.select_dtypes(include='object').columns.tolist()
numerical_features = df.select_dtypes(include=['int64', 'float64']).drop('price', axis=1).columns.tolist()
```

```

if 'bedrooms' in df.columns and 'bathrooms' in df.columns and 'area' in df.columns:
    df['total_rooms'] = df['bedrooms'] + df['bathrooms']
    df['area_per_room'] = df['area'] / df['total_rooms'].replace(0, np.nan)

X = df.drop('price', axis=1)
y = df['price']
numerical_pipeline = Pipeline([
    ('scaler', StandardScaler())
])

categorical_pipeline = Pipeline([
    ('encoder', OneHotEncoder(handle_unknown='ignore', sparse_output=False))
])

preprocessor = ColumnTransformer(transformers=[
    ('num', numerical_pipeline, numerical_features),
    ('cat', categorical_pipeline, categorical_features)
])

X_preprocessed = preprocessor.fit_transform(X)

encoded_cols = preprocessor.named_transformers_['cat']['encoder'].get_feature_names_out(categorical_features)
all_feature_names = numerical_features + list(encoded_cols)

X_train, X_test, y_train, y_test = train_test_split(X_preprocessed, y, test_size=0.2, random_state=42)

print(f"Preprocessed training shape: {X_train.shape}")
print(f"Features used: {all_feature_names}")
preprocessor = ColumnTransformer(transformers=[
    ('num', numerical_pipeline, numerical_features),
    ('cat', categorical_pipeline, categorical_features)
])

X_preprocessed = preprocessor.fit_transform(X)

encoded_cols = preprocessor.named_transformers_['cat']['encoder'].get_feature_names_out(categorical_features)
all_feature_names = numerical_features + list(encoded_cols)

X_train, X_test, y_train, y_test = train_test_split(X_preprocessed, y, test_size=0.2, random_state=42)

print(f"Preprocessed training shape: {X_train.shape}")
print(f"Features used: {all_feature_names}")

```

```

↩
0  13300000  7420    4    2    3    yes    no    no
1  12250000  8960    4    4    4    yes    no    no
2  12250000  9960    3    2    2    yes    no    yes
3  12215000  7500    4    2    2    yes    no    yes
4  11410000  7420    4    1    2    yes    yes   yes

```

```

    hotwaterheating airconditioning parking prefarea furnishingstatus
0                no                yes    2    yes    furnished
1                no                yes    3    no    furnished
2                no                no    2    yes  semi-furnished
3                no                yes    3    yes    furnished
4                no                yes    2    no    furnished

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 545 non-null   int64
1   area                 545 non-null   int64
2   bedrooms             545 non-null   int64
3   bathrooms            545 non-null   int64

```

```

4 stories 545 non-null int64
5 mainroad 545 non-null object
6 guestroom 545 non-null object
7 basement 545 non-null object
8 hotwaterheating 545 non-null object
9 airconditioning 545 non-null object
10 parking 545 non-null int64
11 prefarea 545 non-null object
12 furnishingstatus 545 non-null object

```

dtypes: int64(6), object(7)

memory usage: 55.5+ KB

None

Preprocessed training shape: (436, 20)

Features used: ['area', 'bedrooms', 'bathrooms', 'stories', 'parking', 'mainroad_no', 'mainroad_yes', 'guestro

Preprocessed training shape: (436, 20)

Features used: ['area', 'bedrooms', 'bathrooms', 'stories', 'parking', 'mainroad_no', 'mainroad_yes', 'guestro

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

```

```
df = pd.read_csv('/content/Housing.csv')
```

```
print(df.head())
```

```

X = df.drop('price', axis=1)
y = df['price']

```

```

categorical_cols = X.select_dtypes(include='object').columns.tolist()
numerical_cols = X.select_dtypes(include=[int64, float64]).columns.tolist()

```

```

preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_cols)
    ],
    remainder='passthrough'
)

```

```

model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('regressor', LinearRegression())
])

```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
model.fit(X_train, y_train)
```

```
predictions = model.predict(X_test)
```

```

print("Sample predictions:", predictions[:5])
print("Actual values      :", y_test.values[:5])

```

```

↗
0  13300000  7420  4  2  3  yes  no  no
1  12250000  8960  4  4  4  yes  no  no
2  12250000  9960  3  2  2  yes  no  yes
3  12215000  7500  4  2  2  yes  no  yes
4  11410000  7420  4  1  2  yes  yes  yes

hotwaterheating  airconditioning  parking  prefarea  furnishingstatus
0                no                yes        2        yes        furnished

```

1	no	yes	3	no	furnished
2	no	no	2	yes	semi-furnished
3	no	yes	3	yes	furnished
4	no	yes	2	no	furnished

Sample predictions: [5164653.90033958 7224722.29802165 3109863.24240343 4612075.32722563 3294646.25725961]

Actual values : [4060000 6650000 3710000 6440000 2800000]

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df=pd.read_csv('/content/Housing.csv')
```

```
X = data.drop('price', axis=1)
y = data['price']
X= pd.get_dummies(X, drop_first=True)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

```
↳ /usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge. Increase the number of iterations.
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
y_pred = model.predict(X_test)
print(y_pred,"y_prediction")
```

3500000	3500000	4200000	4200000	4200000	3500000	3500000	4200000	4200000
4200000	3500000	3500000	4200000	4200000	3500000	4200000	4200000	4200000
3500000	3500000	4200000	3500000	3500000	4200000	3500000	3500000	4200000
4200000	3500000	4200000	4200000	4200000	4200000	3500000	3500000	4200000
4200000	4200000	4200000	4200000	3500000	4200000	3500000	4200000	3500000
3500000	4200000	4200000	3500000	4200000	3500000	4200000	3500000	4200000
4200000	4200000	4200000	4200000	3500000	4200000	4200000	3430000	3500000
3500000	3500000	4200000	3500000	4200000	3500000	4200000	4200000	3500000
3500000	3500000	4200000	4200000	4200000	4200000	4200000	3500000	3500000
4200000	4200000	3500000	4200000	4200000	4200000	4200000	4200000	4200000
5495000	3500000	3500000	4200000	3500000	4200000	3500000	3500000	4200000
3500000	4200000	4200000	4200000	4200000	4200000	3500000	4200000	4200000
3500000								

```
from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
model.fit(X_train, y_train)
```

```

    ▾ RandomForestClassifier ⓘ ?
    RandomForestClassifier()

```



```

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
df=pd.read_csv('/content/Housing.csv')
print(accuracy_score(y_test,y_pred))
print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))

```

```

0.009174311926605505
precision    recall  f1-score   support

1750000      0.00      0.00      0.00         1
1820000      0.00      0.00      0.00         1
1855000      0.00      0.00      0.00         0
1890000      0.00      0.00      0.00         2
2100000      0.00      0.00      0.00         1
2233000      0.00      0.00      0.00         1
2275000      0.00      0.00      0.00         1
2380000      0.00      0.00      0.00         1
2450000      0.00      0.00      0.00         2
2520000      0.00      0.00      0.00         1
2653000      0.00      0.00      0.00         0
2660000      0.00      0.00      0.00         4
2695000      0.00      0.00      0.00         0
2800000      0.00      0.00      0.00         1
2870000      0.00      0.00      0.00         1
2940000      0.00      0.00      0.00         2
2975000      0.00      0.00      0.00         0
3003000      0.00      0.00      0.00         1
3010000      0.50      1.00      0.67         1
3045000      0.00      0.00      0.00         1
3080000      0.00      0.00      0.00         2
3115000      0.00      0.00      0.00         0
3150000      0.00      0.00      0.00         1
3220000      0.00      0.00      0.00         1
3234000      0.00      0.00      0.00         1
3290000      0.00      0.00      0.00         1
3325000      0.00      0.00      0.00         1
3353000      0.00      0.00      0.00         1
3360000      0.00      0.00      0.00         2
3430000      0.00      0.00      0.00         0
3465000      0.00      0.00      0.00         0
3500000      0.00      0.00      0.00         3
3640000      0.00      0.00      0.00         1
3675000      0.00      0.00      0.00         1
3703000      0.00      0.00      0.00         1
3710000      0.00      0.00      0.00         1
3773000      0.00      0.00      0.00         1
3780000      0.00      0.00      0.00         1
3836000      0.00      0.00      0.00         0
3850000      0.00      0.00      0.00         1
3885000      0.00      0.00      0.00         0
3920000      0.00      0.00      0.00         0
3990000      0.00      0.00      0.00         0
4007500      0.00      0.00      0.00         1
4060000      0.00      0.00      0.00         1
4165000      0.00      0.00      0.00         1
4193000      0.00      0.00      0.00         2
4200000      0.00      0.00      0.00         1
4270000      0.00      0.00      0.00         1
4340000      0.00      0.00      0.00         2
4473000      0.00      0.00      0.00         0
4480000      0.00      0.00      0.00         1
4543000      0.00      0.00      0.00         1
4550000      0.00      0.00      0.00         2

```

```

from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
print("MAE:", mean_absolute_error(y_test, predictions))
print("MSE:", mean_squared_error(y_test, predictions))
print("R² Score:", r2_score(y_test, predictions))

```

```


MAE: 970043.4039201646
MSE: 1754318687330.7036
R² Score: 0.6529242642153106

```

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(9, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='YlGnBu', linewidths=0.5)
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



```
import numpy as np
import seaborn as sns

sns.heatmap(confusion_matrix(y_test,y_pred),annot=True,fmt='d', cmap='blue')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(5, 4))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation')
plt.show()
```

